

Preliminary Sample-Size Analysis for CLAPS

Pablo Bernabeu and CLAPS collaborators

30 September 2026

Abstract

CLAPS asks whether affectedness, how strongly an event affects the person or thing it happens to, matters more for the acceptability of passive sentences than of active ones, and whether that pattern holds across languages. This report evaluates by simulation what the planned sample of 80 participants per language can be expected to deliver. Studies were simulated from the planned ordinal model fitted to the pilot data in English, Turkish and Norwegian, with the effects drawn from each language’s pilot posterior, and each was analysed as the real data will be. The design has about 72 verbs per language, with agent gender counterbalanced across participants. At 80 participants, the primary prediction, that the affectedness slope is steeper for passives than for actives, reaches a directional Bayes factor of 10 in 90% of simulated English studies and 97% of Norwegian ones. Turkish, whose pilot effect is weakest, reaches it in 57%. The secondary prediction, that the slope is also steeper for passives than for pseudo-passives, reaches a Bayes factor of 10 in 95% of English studies, while the Turkish pilot shows the reverse pattern and the simulations inherit it. Affectedness is a property of the verb, so the verbs bound the precision of every prediction, and beyond 80 participants the rates barely change. The pooled cross-linguistic analysis, planned as a synthesis with no confirmatory claim, reaches a Bayes factor of 10 in 57% of simulated studies with the three pilot languages, at 50 to 100 participants per language. If the seven languages measured so far take part and further languages differ as much as those seven, it does so in 83% of studies with ten languages and 89% with twenty. These rates assume the canonical passive in Mandarin, Indonesian and Balinese, where the choice of passive is still open.

Background

CLAPS (A Crosslinguistic Approach to Passive Semantics) asks whether the acceptability of passive and related sentence types tracks the semantic affectedness of the patient, whether that relationship differs by sentence type and whether it holds across languages. Affectedness, the degree to which a patient or theme undergoes a change of state or is causally impinged upon by the action, is a core dimension of argument structure (Beavers, 2011; Dowty, 1991). It has long been argued to be the semantic heart of the passive, which carries the meaning that the surface subject is in a state or circumstance characterised by the agent having acted upon it (Pinker et al., 1987). Acceptability-judgement and comprehension studies with adults in several languages support the view that passive acceptability is semantically structured (Ambridge et al., 2016; Ambridge et al., 2023; Aryawibawa & Ambridge, 2018; Bidgood et al., 2020; Darmasetiyawan & Ambridge, 2022; Liu & Ambridge, 2021).

The evidence to date spans English, Indonesian, Balinese, Hebrew and Mandarin. A Bayesian meta-analytic synthesis of those studies finds that affectedness raises the acceptability of passives but not of pseudo-passives (Ambridge et al., 2023). In most of the five languages, it raises passive acceptability more than active acceptability, although the size of that advantage varies considerably and in Balinese it runs the other way. CLAPS adds Turkish and Norwegian, for which its pilot supplies estimates.

Each language’s data will be collected by its own team, at a sample size fixed in advance. This report, prepared for the study’s Stage 1 Registered Report, evaluates what the planned sample can be expected to deliver. It does so through a design analysis (Schönbrodt & Wagenmakers, 2018) whose simulated effects come from the CLAPS pilot.

The Model and the Predictions

Each acceptability judgement is a rating from 1 to 7. The ratings are modelled with a Bayesian cumulative-logit mixed-effects model, an ordinal model that respects the discreteness, ordering and bounded range of rating data (Bürkner & Vuorre, 2019). Treating the same ratings as continuous biases the regression coefficients and inflates error rates where floor or ceiling effects are present (Liddell & Kruschke, 2018), and modelling them cumulatively recovers psycholinguistic effects more accurately (Taylor et al., 2023; Veríssimo, 2021).


The model predicts acceptability from sentence type, affectedness and their interaction. Each stimulus appears as an active, a passive and a pseudo-passive. In English, for example, these are “The woman pushed the man”, “The man was pushed by the woman” and “It was the man that the

woman pushed”, the last a cleft that puts the patient first without the passive’s morphology. The passive is the reference level, and the active and the pseudo-passive are compared with it.

The pilot, whose sample Appendix A describes, shows the expected pattern in English and Norwegian (Figure 1). In both, the affectedness slope is positive and steepest for the passive, and flat or slightly negative for the active, as it is for the English pseudo-passive. In Turkish, the passive slope is also steeper than the active slope, though by less, while the pseudo-passive slope is at least as steep as the passive slope. The trends look much the same for male and female agents, although in English the active-by-affectedness interaction differs between agent genders. The design therefore balances the two agent genders across participants (Appendix B).

Figure 1

Raw Pilot Acceptability by Affectedness, Sentence Type, Agent Gender and Language



Note. Panels are by language across rows and by agent gender across columns. Points are individual pilot ratings, jittered, and lines are per-sentence-type linear trends with 95% confidence bands, which ignore the clustering of ratings within participants and verbs and so understate the uncertainty. Affectedness is the score normed for each agent-gender version, z-scored within language, and the model uses one score per verb, the mean of the two. The Norwegian materials include no pseudo-passive.

H1a is the passive affectedness slope, the effect of affectedness on the acceptability of passives. It is a planned estimate, reported with its 95% credible interval but with no prediction or test of its own. The theory’s claim concerns the passive in particular, and affectedness could raise the acceptability of sentences in general, so the tests compare the passive slope with the active and pseudo-passive slopes. H1b, the primary prediction, holds that the affectedness slope is steeper for passives than for actives, so that the active-by-affectedness interaction is negative. H2b, the secondary prediction, holds that the slope is also steeper for passives than for pseudo-passives, so that the pseudo-passive-by-affectedness interaction is negative. The meta-analysis ([Ambridge et al., 2023](#)) and the English pilot both put that interaction below zero. H2a, the opposite direction, is not proposed as a hypothesis, although a Bayes factor for H2b of 1/10 or less is evidence for it. H1b is tested in every language, and H2b in every language with a pseudo-passive, which the Norwegian materials do not include.

Design-Analysis Method

The design analysis simulates many studies from the cumulative-logit model planned for the real data and analyses each as the real data will be analysed. The procedure stays within the Bayesian framework throughout, because the confirmatory test is a Bayes factor. Its result is the probability that a study reaches a Bayes factor of 10, a different quantity from the power to return a significant p -value.

The per-language simulations take their effects from the pilot. For each language, the planned model is fitted to the pilot ratings. On every simulated study, the fixed effects and ordinal thresholds are drawn together from one draw of that fit’s posterior, and the random-effect covariance is held at its posterior mean. New participants and new verb deviations are then drawn around the pilot’s own verbs and their measured affectedness scores. Each reported rate is an assurance, a detection rate averaged over the effect sizes the pilot leaves plausible, which allows for the uncertainty of a single pilot’s point estimate ([Albers & Lakens, 2018](#)). Every pilot language goes forward whatever its pilot showed, so the follow-up bias that Albers and Lakens describe does not arise.

Published effect sizes are subject to publication bias and tend to overstate true effects ([Button et al., 2013](#); [Gelman & Carlin, 2014](#)), so design analyses anchored to them can be optimistic ([Vasishth et al., 2018](#)). In this report, published estimates enter only the multi-language synthesis, for the four languages the pilot does not cover, and there they carry extra uncertainty to allow for the change of protocol. Their average, -0.34 per standard deviation of affectedness, is close to the three pilots’ -0.29 , so they show no sign of inflation relative to the CLAPS protocol.

Affectedness is a property of the verb. In the study, it will be rated in a separate norming task, in which twenty further speakers of each language rate the affectedness of the patient in each event,

once for each agent-gender version, and a verb’s score is the mean of its two version scores. Each verb therefore carries one affectedness score, the same across participants and sentence types, and every affectedness effect is estimated from differences between verbs. The verbs and the spread of their scores bound the precision of those effects, however many participants take part. This is the stimulus-sampling limit long recognised in psycholinguistics (Clark, 1973) and quantified for designs with crossed random effects by Westfall et al. (2014; see also Brysbaert & Stevens, 2018; Judd et al., 2017). With unlimited participants, the standard error of such an effect cannot fall below the by-verb standard deviation of that effect divided by the square root of the verbs’ summed squared affectedness deviations, a limit that can be read from the fitted pilot models without simulation.

The analysis model has the maximal random-effects structure the design supports (Barr et al., 2013), with correlated by-participant slopes for sentence type, affectedness and their interaction, and by-verb intercepts and sentence-type slopes. Because affectedness does not vary within a verb, a by-verb affectedness slope cannot be estimated and is omitted. Affectedness is centred and divided by twice its standard deviation within each language, and effects are reported per standard deviation of affectedness. The priors are centred at zero on the focal slopes, so they favour neither sign. The affectedness slope and the active interaction take normal priors with a standard deviation of 0.5 on the model’s scale, or 0.25 per standard deviation of affectedness, and the pseudo-passive interaction takes 0.6, or 0.30 per standard deviation. The other fixed effects take normal priors with a standard deviation of 1.5. The ordinal thresholds take Student-*t* priors with three degrees of freedom and a scale of 2.5, and the random-effect standard deviations Student-*t* priors with three degrees of freedom and a scale of 1. Their correlations take an LKJ prior with a shape of 2 (Gelman, 2006; Lewandowski et al., 2009).

Evidence for each tested prediction is a directional Bayes factor, the posterior odds that the coefficient carries the predicted sign divided by the same odds under the prior (Morey & Wagenmakers, 2014). It is estimated from posterior draws and from prior draws sampled alongside them. The two signs it compares are each represented by the zero-centred prior restricted to that side, which is a half-normal. Under these symmetric priors, a Bayes factor of 10 corresponds to a posterior probability of .91 that the effect has the predicted sign, and a Bayes factor of 1/10 or less is the same weight of evidence for the opposite sign. A simulated study counts as a detection when its Bayes factor reaches 10.

Each stimulus exists in two versions, one with a woman agent and one with a man agent. Each participant will see one version of each verb, so an English or Turkish session has 216 sentences, and the simulations describe that session. The versions are balanced across participants by two complementary lists (Appendix B).

The per-language results rest on a dedicated set of 120 simulated studies at each of 70, 80 and 100 participants per language, each scored for every prediction, at the pilot’s verb counts of 72 in English and Turkish and 66 in Norwegian. The repository’s wider grid adds 24 earlier studies at 70 and 100 participants, scored for H1b alone. Each study is fitted with four chains of 3,000 iterations

at a target acceptance rate of .99 and a maximum treedepth of 12. After 1,000 warm-up iterations per chain, 8,000 post-warm-up draws remain. The confirmatory analysis will use sixteen chains of 5,000 iterations with 2,000 of warm-up. The shorter chains add Monte-Carlo noise to each simulated Bayes factor without shifting it systematically. Each per-language rate in Table 1 carries an exact 95% interval. Over 120 studies, a rate near 90% has a Monte-Carlo standard error of about three percentage points, so differences of a few points between neighbouring sample sizes are noise.

The simulations ran on the University of Oxford’s Advanced Research Computing service, under R 4.4.2 with brms 2.21.0 and the cmdstanr backend on CmdStan 2.39.0, and the lockfile `renv.lock` fixes the package versions. The analysis code, the simulation settings and the aggregated results are openly available at <https://github.com/pablobernabeu/CLAPS-analysis>, and this report rebuilds from those results without the raw pilot data, which are not yet public.

Results

Table 1 gives the detection rate for each prediction by language and sample size.

Table 1

Detection Rates by Language and Participant Count

Language	Participants	H1b	H2b	H2b reversed
English	70	79% [71, 86]	93% [87, 97]	0% [0, 3]
	80	90% [83, 95]	95% [89, 98]	0% [0, 3]
	100	89% [82, 94]	99% [95, 100]	0% [0, 3]
Turkish	70	56% [46, 65]	1% [0, 5]	72% [63, 80]
	80	57% [47, 66]	0% [0, 3]	82% [74, 88]
	100	54% [45, 63]	0% [0, 3]	77% [68, 84]
Norwegian	70	95% [89, 98]	—	—
	80	97% [92, 99]	—	—
	100	98% [93, 99]	—	—

Note. Each row rests on 120 simulated studies. Entries are the percentage of studies reaching a directional Bayes factor of 10, with an exact 95% interval in brackets. H1b, the primary prediction, is the active-by-affectedness interaction, and H2b, the secondary, is the pseudo-passive-by-affectedness interaction, both predicted negative. The last column gives the percentage of studies in which H2b’s Bayes factor is 1/10 or less, evidence for the opposite direction. The Norwegian materials include no pseudo-passive.

The Primary Prediction

At 80 participants, H1b reaches a Bayes factor of 10 in 90% [83, 95] of simulated English studies, 97% [92, 99] of Norwegian ones and 57% [47, 66] of Turkish ones. The Turkish pilot interaction is the weakest of the three, at -0.15 logits per standard deviation of affectedness against -0.44 in English and -0.27 in Norwegian.

The rates change little between 80 and 100 participants, as the verb limit predicts, so recruiting beyond 80 would add little in any language. At that limit, averaged over the pilot posterior, the interaction would reach a Bayes factor of 10 in about 97% of English studies, over 99% of Norwegian ones and about 65% of Turkish ones.

The Secondary Prediction

In English, the pilot's pseudo-passive-by-affectedness interaction is -0.38 logits per standard deviation of affectedness, with a 95% credible interval from -0.56 to -0.20 and all 8,000 posterior draws below zero. At 80 participants, H2b reaches a Bayes factor of 10 in 95% [89, 98] of simulated studies, more often than H1b.

In Turkish, the pilot interaction runs the other way, at $+0.16$ per standard deviation, with a 95% credible interval from $+0.03$ to $+0.28$, so the affectedness slope is steeper for Turkish pseudo-passives than for passives. The simulations draw their effects from that posterior and inherit the reversal. At 80 participants, 82% of simulated Turkish studies give a Bayes factor for H2b of 1/10 or less, and the remaining 22 of 120 are inconclusive. A reversal of this kind would be a finding in its own right, showing that the constraint on pseudo-passives differs between languages. The Stage 1 protocol will state in advance how it will be interpreted.

The negative pseudo-passive interaction also appears in Balinese, Indonesian and Mandarin in the per-language models of the meta-analysis, although each sign depends on which construction counts as the passive in that language, a choice taken up in the next section. Those studies inform the planning, but the confirmatory data will be collected under the CLAPS protocol alone.

The Pooled Synthesis

The pooled model tests the primary prediction, the active-by-affectedness interaction, across all the languages together. It has the per-language random-effects structure with random slopes by language added, so its population-level interaction estimates the cross-linguistic effect and each language can depart from it. It is planned as a synthesis with no confirmatory claim, reported at whatever number of languages takes part, and the confirmatory claims rest on the per-language

tests. Pooled figures for the pseudo-passive interaction await the ruling on the passive in Mandarin, Indonesian and Balinese.

With the three pilot languages alone, the pooled model reaches a Bayes factor of 10 on the interaction in 57% [47, 67] of 107 simulated studies at 50 to 100 participants per language. The rate shows no trend over that range. It is held down by the variation of the effect between languages, which neither verbs nor participants reduce and which three languages estimate poorly. These simulations used the structure of the published meta-analysis model, which also carries a by-verb affectedness slope.

Because more teams are expected to take part, the synthesis is also simulated with ten and twenty languages. A two-stage version simulates each language’s estimate of the interaction and combines the estimates in a Bayesian random-effects meta-analysis. Its one free parameter, the prior that the full model in effect places on the between-language spread, is calibrated against the full model’s three-language simulations. In these simulations, the seven languages measured so far take part as CLAPS samples, with effects drawn from their estimates. They are the three CLAPS pilots, and Balinese, Hebrew, Indonesian and Mandarin, the four languages of the meta-analysis other than English. The remaining languages, up to ten or twenty, are drawn from a population of languages, and Table 2 gives the result.

Table 2

Pooled Detection Rate with the Seven Measured Languages and Further Languages Drawn from a Population

New languages as different as	Median τ	Ten languages	Twenty languages	Twenty, alternative reading
The three CLAPS pilots	0.16	77%	92%	69%
All seven measured languages	0.39	83%	89%	23%
The four other meta-analysis languages	0.53	73%	76%	12%

Note. Each rate is the percentage of 4,000 simulated studies in which the pooled directional Bayes factor on the active-by-affectedness interaction reaches 10, with 80 participants and 72 verbs per language. τ is the between-language standard deviation of the interaction, per standard deviation of affectedness, and the second column gives its posterior median for the set in the first column. New languages are drawn from a population whose mean and τ are themselves drawn from that posterior. The seven measured languages take part at effects drawn from their estimates. Each of the four from the meta-analysis carries an extra standard deviation, set at 0.15 to allow for the change of protocol. In the first row it is set at 0.40, since a spread as small as the pilots’ implies that the earlier estimates carry over poorly. Holding the seven at their point estimates moves no rate in the second and third rows by more than 5 points. The first row is left out of that comparison, since

holding the seven fixed contradicts its premise. The last column takes the other construction as the passive in Mandarin, Indonesian and Balinese, which puts τ at 0.52 for all seven languages and 0.62 for the four.

The spread of the effect between languages decides the result, and it is not yet known. The three CLAPS pilots differ little, the four other meta-analysis languages differ a great deal, and all seven together lie between. The larger spread may reflect real differences between the languages, or the earlier studies' procedures, which differed in scale length, verb count, affectedness norming and the construction treated as the passive. The spread across all seven languages is the planning value used here, since three languages are too few to estimate a spread well and the larger spread may be partly procedural. On that value, the pooled test reaches a Bayes factor of 10 in 83% of simulated studies with ten languages and in 89% with twenty. If new languages differ as little as the pilots, the rate falls to 77% at ten languages, because the four meta-analysis estimates then carry over poorly, and rises to 92% at twenty.

The result also turns on which construction counts as the passive in Mandarin, Indonesian and Balinese, each of which has more than one passive-like construction. The figures above take the canonical passive in each, as the meta-analysis did. Taking the other construction as the passive instead changes the sign of the interaction in Indonesian and Mandarin and moves the average across the seven languages from -0.32 to -0.03 . With twenty languages, the pooled test then reaches a Bayes factor of 10 in only 23% of studies on the seven-language spread. Which reading applies is a linguistic decision, to be made for each language before its team collects data.

The Sampling Plan

Every team will collect usable data from 80 participants, replacing any participant who is excluded, a number fixed in advance because each team must state its sample size to its ethics committee. Each language will use as close to 72 verbs as its supply of suitable verbs allows.

Nature Human Behaviour sets its 0.95 power requirement for frequentist analysis plans. For a Bayesian plan, it asks that data collection continue until the Bayes factor is “at least 10 times in favour of the experimental hypothesis over the null hypothesis (or vice versa)”. Authors with limited resources may instead state a maximum feasible sample, “sufficiently large that inconclusive results at this sample size would nevertheless be an important message for the field”. CLAPS would take the second route for every hypothesis, with a Bayes factor of 10 as the threshold in every language. Eighty is the most the collecting teams can provide. At that sample, each language is close to the limit its verbs set, so further participants would add little to the chance of reaching the threshold. An inconclusive result at 80 would still reflect an estimate with nearly the precision the language's verbs allow.

The Bayes factor to be registered is directional ([Morey & Wagenmakers, 2014](#)). Each prediction is represented by the half-normal given in the Design-Analysis Method, one of the forms the journal names for representing a prediction, and its competitor is the reversed effect, represented by the same half-normal on the other side. A Bayes factor of 10 is therefore strong evidence for the prediction over its reversal, and one of 1/10 or less is equally strong evidence the other way.

Turkish would be registered by the same route. Its interaction reaches a Bayes factor of 10 in 57% of simulated studies and in at most about 65% with unlimited participants, so an inconclusive Turkish result would be informative for the reason given above. Its pseudo-passive interaction will probably run the other way. Both Turkish predictions would be registered as hypotheses, because the theory makes the same predictions for every language, and leaving a language out on the strength of its pilot would amount to choosing languages by their results. [Table 3](#) summarises the planned analyses and their detection rates.

Table 3*Planned Analyses and Their Detection Rates*

Analysis	Coefficient and direction	Test	Detection rate
H1b, primary	Active by affectedness, negative	Directional Bayes factor of 10, in each language	English 90%, Norwegian 97% and Turkish 57%
H2b, secondary	Pseudo-passive by affectedness, negative	Directional Bayes factor of 10, in each language with a pseudo-passive	English 95%. Turkish is expected to reverse, with a Bayes factor of 1/10 or less in 82%
H1a	Passive affectedness slope	Planned estimate, with its 95% credible interval	Not tested
Pooled synthesis	Cross-linguistic active-by-affectedness interaction, negative	Reported with its directional Bayes factor, as a synthesis with no confirmatory claim	57% with the three pilot languages. Including the seven measured languages, 83% with ten and 89% with twenty

Note. The per-language rates are at 80 participants. The three-language pooled rate combines simulations at 50 to 100 participants per language. The ten- and twenty-language rates are at 80 participants per language, on the spread across the seven measured languages and the canonical reading of the passive.

Limitations and Open Decisions

At 80 participants, no simulated fit has a divergent transition or a saturated treedepth. Confirmatory fits must also reach an R -hat below 1.01 and bulk and tail effective sample sizes of at least 400. With the shorter runs used for simulation, 27 of 120 English fits meet those thresholds, as do 33 Norwegian and 12 Turkish fits. Every rate counts all fits. Over the passing fits alone, the rates are 81% in English, 100% in Norwegian and 50% in Turkish, each within what a random subset of that size would show. The confirmatory chains will also be half as long again, and sixteen of them will give six times the post-warm-up draws.

Fitted to the combined pilot data, 75,528 ratings over 210 verbs, the pooled model estimated the interaction at about -0.24 logits per standard deviation of affectedness, with a 95% credible interval from -0.51 to $+0.15$. That estimate is close to the average of the three pilot estimates, -0.29 . The pooled model is harder to fit than the per-language ones, however. It saturated the maximum

treedepth on three quarters of its iterations, and confirmatory fits must show none. A pooled fit to the collected data that misses the criteria falls back to the two-stage synthesis of the per-language estimates, the analysis behind the ten- and twenty-language rates above.

The multi-language figures rest on an assumed population of languages. The seven measured so far are a convenience sample from five families in two macroareas, and which further languages will join is not yet known.

Several decisions remain open. Mandarin, Indonesian and Balinese each need a ruling on which construction is the passive, and Turkish needs confirmation that its pseudo-passive is the same construction as the English one. It has yet to be settled which teams will collect data, and whether the Stage 2 paper waits for a minimum number of languages or the pooled synthesis is reported whatever the number. The pooled model also has to specify whether a verb and its translations share one by-verb random effect. The maximum-feasible route for every hypothesis, the treatment of H1a as a planned estimate, registering only the negative direction of H2b and the registration of both Turkish predictions as hypotheses await the collaborators' agreement. None of these decisions changes the per-language sample size.

References

- Albers, C., & Lakens, D. (2018). When power analyses based on pilot data are biased: Inaccurate effect size estimators and follow-up bias. *Journal of Experimental Social Psychology*, 74, 187–195. <https://doi.org/10.1016/j.jesp.2017.09.004>
- Ambridge, B., Arnon, I., & Bekman, D. (2023). He was run-over by a bus: Passive – but not pseudo-passive – sentences are rated as more acceptable when the subject is highly affected. New data from Hebrew, and a meta-analytic synthesis across English, Balinese, Hebrew, Indonesian and Mandarin. *Glossa Psycholinguistics*, 2(1), 11. <https://doi.org/10.5070/G6011177>
- Ambridge, B., Bidgood, A., Pine, J. M., Rowland, C. F., & Freudenthal, D. (2016). Is passive syntax semantically constrained? Evidence from adult grammaticality judgment and comprehension studies. *Cognitive Science*, 40(6), 1435–1459. <https://doi.org/10.1111/cogs.12277>
- Aryawibawa, I. N., & Ambridge, B. (2018). Is syntax semantically constrained? Evidence from a grammaticality judgment study of Indonesian. *Cognitive Science*, 42(8), 3135–3148. <https://doi.org/10.1111/cogs.12697>
- Barr, D. J., Levy, R., Scheepers, C., & Tily, H. J. (2013). Random effects structure for confirmatory hypothesis testing: Keep it maximal. *Journal of Memory and Language*, 68(3), 255–278. <https://doi.org/10.1016/j.jml.2012.11.001>
- Beavers, J. (2011). On affectedness. *Natural Language & Linguistic Theory*, 29(2), 335–370. <https://doi.org/10.1007/s11049-011-9124-6>
- Bidgood, A., Pine, J. M., Rowland, C. F., & Ambridge, B. (2020). Syntactic representations are both abstract and semantically constrained: Evidence from children’s and adults’ comprehension and production/priming of the English passive. *Cognitive Science*, 44(9), e12892. <https://doi.org/10.1111/cogs.12892>
- Brysbaert, M., & Stevens, M. (2018). Power analysis and effect size in mixed effects models: A tutorial. *Journal of Cognition*, 1(1), 9. <https://doi.org/10.5334/joc.10>
- Bürkner, P.-C., & Vuorre, M. (2019). Ordinal regression models in psychology: A tutorial. *Advances in Methods and Practices in Psychological Science*, 2(1), 77–101. <https://doi.org/10.1177/2515245918823199>
- Button, K. S., Ioannidis, J. P. A., Mokrysz, C., Nosek, B. A., Flint, J., Robinson, E. S. J., & Munafò, M. R. (2013). Power failure: Why small sample size undermines the reliability of neuroscience. *Nature Reviews Neuroscience*, 14(5), 365–376. <https://doi.org/10.1038/nrn3475>
- Clark, H. H. (1973). The language-as-fixed-effect fallacy: A critique of language statistics in psychological research. *Journal of Verbal Learning and Verbal Behavior*, 12(4), 335–359. [https://doi.org/10.1016/S0022-5371\(73\)80014-3](https://doi.org/10.1016/S0022-5371(73)80014-3)
- Darmasetiyawan, I. M. S., & Ambridge, B. (2022). Syntactic representations contain semantic information: Evidence from Balinese passives. *Collabra: Psychology*, 8(1), 33133. <https://doi.org/10.1525/collabra.33133>

- Dowty, D. R. (1991). Thematic proto-roles and argument selection. *Language*, 67(3), 547–619. <https://doi.org/10.1353/lan.1991.0021>
- Gelman, A. (2006). Prior distributions for variance parameters in hierarchical models (comment on article by Browne and Draper). *Bayesian Analysis*, 1(3), 515–534. <https://doi.org/10.1214/06-BA117A>
- Gelman, A., & Carlin, J. (2014). Beyond power calculations: Assessing Type S (sign) and Type M (magnitude) errors. *Perspectives on Psychological Science*, 9(6), 641–651. <https://doi.org/10.1177/1745691614551642>
- Judd, C. M., Westfall, J., & Kenny, D. A. (2017). Experiments with more than one random factor: Designs, analytic models, and statistical power. *Annual Review of Psychology*, 68, 601–625. <https://doi.org/10.1146/annurev-psych-122414-033702>
- Lewandowski, D., Kurowicka, D., & Joe, H. (2009). Generating random correlation matrices based on vines and extended onion method. *Journal of Multivariate Analysis*, 100(9), 1989–2001. <https://doi.org/10.1016/j.jmva.2009.04.008>
- Liddell, T. M., & Kruschke, J. K. (2018). Analyzing ordinal data with metric models: What could possibly go wrong? *Journal of Experimental Social Psychology*, 79, 328–348. <https://doi.org/10.1016/j.jesp.2018.08.009>
- Liu, L., & Ambridge, B. (2021). Balancing information-structure and semantic constraints on construction choice: Building a computational model of passive and passive-like constructions in Mandarin Chinese. *Cognitive Linguistics*, 32(3), 349–388. <https://doi.org/10.1515/cog-2019-0100>
- Morey, R. D., & Wagenmakers, E.-J. (2014). Simple relation between Bayesian order-restricted and point-null hypothesis tests. *Statistics & Probability Letters*, 92, 121–124. <https://doi.org/10.1016/j.spl.2014.05.010>
- Pinker, S., Lebeaux, D. S., & Frost, L. A. (1987). Productivity and constraints in the acquisition of the passive. *Cognition*, 26(3), 195–267. [https://doi.org/10.1016/S0010-0277\(87\)80001-X](https://doi.org/10.1016/S0010-0277(87)80001-X)
- Schönbrodt, F. D., & Wagenmakers, E.-J. (2018). Bayes factor design analysis: Planning for compelling evidence. *Psychonomic Bulletin & Review*, 25(1), 128–142. <https://doi.org/10.3758/s13423-017-1230-y>
- Taylor, J. E., Rousselet, G. A., Scheepers, C., & Sereno, S. C. (2023). Rating norms should be calculated from cumulative link mixed effects models. *Behavior Research Methods*, 55(5), 2175–2196. <https://doi.org/10.3758/s13428-022-01814-7>
- Vasishth, S., Mertzen, D., Jäger, L. A., & Gelman, A. (2018). The statistical significance filter leads to overoptimistic expectations of replicability. *Journal of Memory and Language*, 103, 151–175. <https://doi.org/10.1016/j.jml.2018.07.004>
- Veríssimo, J. (2021). Analysis of rating scales: A pervasive problem in bilingualism research and a solution with Bayesian ordinal models. *Bilingualism: Language and Cognition*, 24(5), 842–848. <https://doi.org/10.1017/S1366728921000316>

Westfall, J., Kenny, D. A., & Judd, C. M. (2014). Statistical power and optimal design in experiments in which samples of participants respond to samples of stimuli. *Journal of Experimental Psychology: General*, 143(5), 2020–2045. <https://doi.org/10.1037/xge0000014>

Appendix A

Composition of the Pilot Sample

This appendix describes the pilot sample whose ratings give the descriptive trends in Figure 1 and to which the planned model is fitted for the simulations. Table A1 gives the counts, read directly from the harmonised pilot data.

Table A1

Composition of the Harmonised Pilot Sample, by Language

Language	Participants	Verbs	Sentence types	Items	Judgements
English	70	72	3	144	30,240
Turkish	70	72	3	144	30,240
Norwegian	57	66	2	132	15,048

Note. Items are verb-by-agent-gender combinations, so the item count is twice the verb count. Every pilot participant saw both versions of every verb, whereas in the study itself each participant will see one version of each verb (Appendix B). The Norwegian materials include no pseudo-passive, hence only two sentence types. Norwegian sessions also included a synthetic passive, which the preregistered analysis excludes, so a Norwegian participant rated one and a half times as many sentences as this table counts.

Appendix B

The Agent-Gender Lists

This appendix sets out how the two agent-gender versions of each stimulus are balanced across participants, and why. Each participant will see one version of each verb, so the balance has to come from the assignment of versions to participants.

Balance matters because agent gender moves the interaction in English. In a linear mixed model of the 1-to-7 pilot ratings, used here only to gauge the relative size of the gender term, with agent gender coded plus or minus a half, the English interaction is -0.16 rating points per standard deviation of affectedness and its gender-by-active-by-affectedness term is $+0.068$ ($SE = 0.023$). The interaction is therefore markedly stronger for man-agent items than for woman-agent ones. A design in which every participant saw only woman-agent versions, or only man-agent ones, would move the English interaction by about 21%, against 1% in Turkish and 3% in Norwegian.

Two complementary lists remove that shift without lengthening the session. The verbs are split into two sets of 36, or 33 in Norwegian, matched on affectedness and verb class, and the participants into two lists. List 1 shows the first set with woman agents and the second with man agents, and list 2 the reverse. Each participant therefore sees both agent genders equally often, each verb is rated in each version by half the sample, and the estimand remains the slope averaged over agent gender. Because both lists show both agent genders on matched verb sets, an unequal split between the lists leaves the interaction essentially unchanged. Even so, assigning participants to the lists alternately, and replacing exclusions within the same list, keeps each verb’s two versions evenly rated. Splitting the agent-gender versions across participants in this way follows Ambridge et al. (2016, Study 3) and the Hebrew study of Ambridge et al. (2023).

Both versions stay in the affectedness rating task, which asks for one judgement per verb and agent gender and is short. Halving it would leave each agent-specific score resting on half as many raters and would attenuate every slope estimated from those scores.