Preliminary Sample-Size Analysis for CLAPS

Pablo Bernabeu and CLAPS collaborators

13 October 2026

Abstract

CLAPS asks whether affectedness, how strongly an event affects the person or thing it happens to, influences the acceptability of passive sentences more than that of active ones, and whether that pattern holds across languages. The primary hypothesis (H1) holds that the affectedness slope is steeper for passives than for actives, and the secondary hypothesis (H2) that it is also steeper for passives than for pseudo-passives. This report evaluates by simulation what the planned design, 80 participants per language in at least 20 languages, can be expected to deliver. Both hypotheses are tested across languages in a two-stage synthesis: each language is analysed with the planned ordinal model, and the estimates are combined in a Bayesian random-effects meta-analysis. The simulations draw the effects of English, Turkish and Norwegian from their pilot posteriors, those of Balinese, Hebrew, Indonesian and Mandarin from published estimates, and those of further languages from a population like these seven. With 20 languages, the synthesis reaches a directional Bayes factor of 10 in 88% of simulated studies on H1, and in 89% on H2 across the 15 languages with a pseudo-passive. These rates depend mainly on how much the effect varies between languages. Within a single pilot language at 80 participants, H1 reaches a Bayes factor of 10 in 90% of simulated English studies, 97% of Norwegian ones and 57% of Turkish ones. H2 reaches it in 95% of English studies, whereas the Turkish pilot points the other way. Because affectedness is a property of the verb, the verbs bound the precision of each language's estimate, and beyond 80 participants the rates barely change.

Background

CLAPS (A Crosslinguistic Approach to Passive Semantics) asks whether the acceptability of passive and related sentence types tracks the semantic affectedness of the patient, whether that relationship differs by sentence type and whether it holds across languages. Affectedness, the degree to which a patient or theme undergoes a change of state or is causally impinged upon by the action, is a core dimension of argument structure (Beavers, 2011; Dowty, 1991). It has long been argued to be the semantic heart of the passive, which carries the meaning that the surface subject is in a state or circumstance characterised by the agent having acted upon it (Pinker et al., 1987). Acceptability-judgement and comprehension studies with adults in several languages support the view that passive acceptability is semantically structured (Ambridge et al., 2016; Ambridge et al., 2023; Aryawibawa & Ambridge, 2018; Bidgood et al., 2020; Darmasetiyawan & Ambridge, 2022; Liu & Ambridge, 2021).

The evidence to date spans English, Indonesian, Balinese, Hebrew and Mandarin. A Bayesian meta-analytic synthesis of those studies finds that affectedness raises the acceptability of passives but not of pseudo-passives (Ambridge et al., 2023). In most of the five languages, it raises passive acceptability more than active acceptability, although the size of that advantage varies considerably and in Balinese it runs the other way. The CLAPS pilot, run in English, Turkish and Norwegian, adds two languages to that evidence.

Each language's data will be collected by its own team, at a sample size fixed in advance. This report, prepared for the study's Stage 1 Registered Report, evaluates what the planned sample can be expected to deliver. It does so through a design analysis (Schönbrodt & Wagenmakers, 2018) whose simulated effects come from the CLAPS pilot and, for the languages the pilot does not cover, from the published estimates.

The Model and the Predictions

Each acceptability judgement is a rating from 1 to 7. The ratings are modelled with a Bayesian cumulative-logit mixed-effects model, an ordinal model that respects the discreteness, ordering and bounded range of rating data (Bürkner & Vuorre, 2019). Treating the same ratings as continuous biases the regression coefficients and inflates error rates where floor or ceiling effects are present (Liddell & Kruschke, 2018), and modelling them cumulatively recovers psycholinguistic effects more accurately (Taylor et al., 2023; Veríssimo, 2021).

The model predicts acceptability from sentence type, affectedness and their interaction. Each stimulus appears as an active, a passive and a pseudo-passive. In English, for example, these are "The woman pushed the man", "The man was pushed by the woman" and "It was the man that the

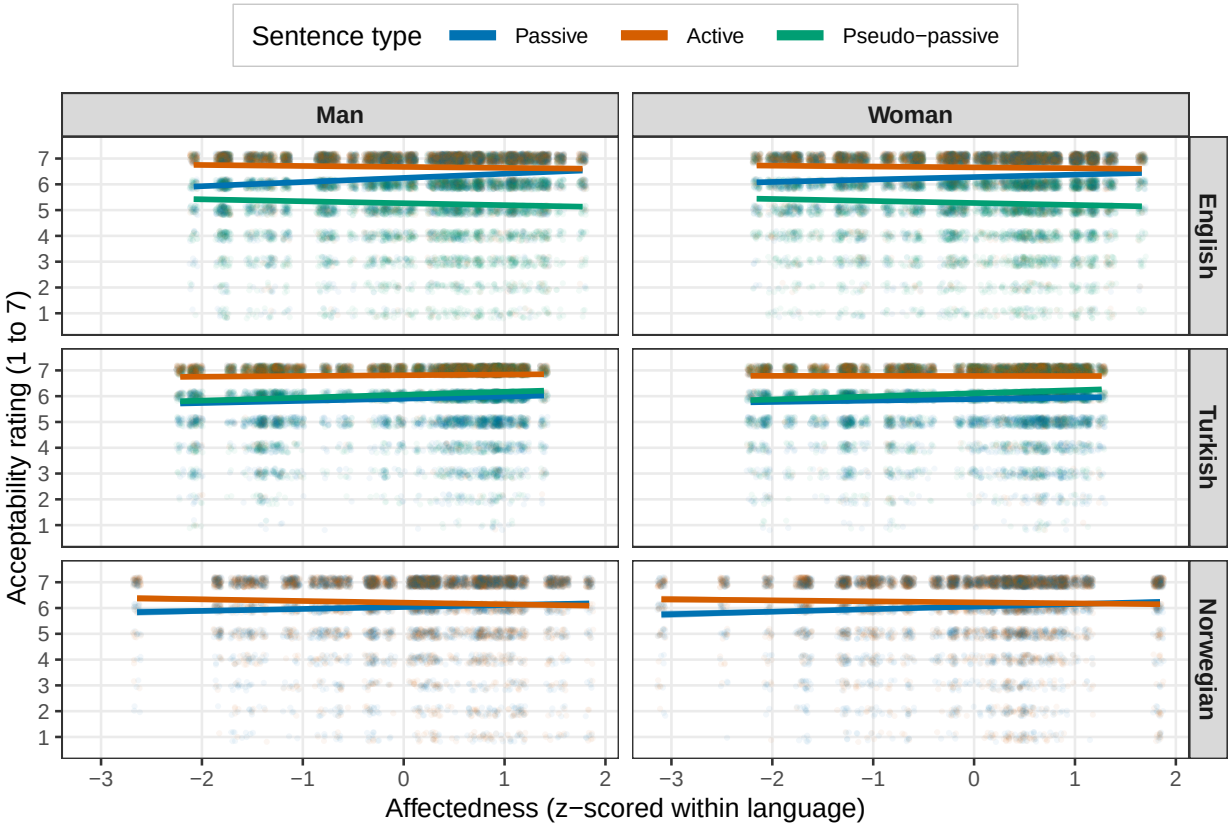
woman pushed", the last a cleft that puts the patient first without the passive's morphology. Each sentence type is treated as the same construction in every language that has it. The passive is the reference level, and the active and the pseudo-passive are compared with it.

H1, the primary hypothesis, holds that the affectedness slope is steeper for passives than for actives, so that the active-by-affectedness interaction is negative. H2, the secondary hypothesis, holds that the slope is also steeper for passives than for pseudo-passives, so that the pseudo-passive-by-affectedness interaction is negative. The passive slope itself carries no prediction. Affectedness might raise or lower the acceptability of sentences in general, and the theory requires only that its slope be steeper for passives, so each test compares the passive slope with another slope, whatever their signs. H1 is tested on the mean across all languages, and H2 on the mean across the languages with a pseudo-passive, which the Norwegian materials do not include.

The pilot, whose sample Appendix A describes, shows the predicted pattern in English and Norwegian (Figure 1). In both, the affectedness slope is positive and steepest for the passive, and flat or slightly negative for the active, as it is for the English pseudo-passive. In Turkish, the passive slope is also steeper than the active slope, though by less, while the pseudo-passive slope is at least as steep as the passive slope. The trends look much the same for male and female agents, although in English the active-by-affectedness interaction differs between them. The design therefore balances the two agent genders across participants (Appendix B).

Figure 1

Raw Pilot Acceptability by Affectedness, Sentence Type, Agent Gender and Language



Note. Panels are by language across rows and by agent gender across columns. Points are individual pilot ratings, jittered. Lines are linear trends for each sentence type, with 95% confidence bands that ignore the clustering of ratings within participants and verbs and so understate the uncertainty. Affectedness is the score normed for each agent-gender version, z-scored within language, and the model uses one score per verb, the mean of the two. The Norwegian materials include no pseudo-passive.

Design-Analysis Method

The design analysis simulates many studies and analyses each one as the real data will be analysed. Simulations of single languages, from the cumulative-logit model planned for each language's data, show how precisely one language's data estimate each interaction. Simulations of 20 languages then evaluate the confirmatory test, which combines the per-language estimates in the way described below. The procedure stays within the Bayesian framework throughout, because the confirmatory

test is a Bayes factor. Its result is therefore the probability that a study reaches a Bayes factor of 10, a different quantity from the power to return a significant p-value.

The per-language simulations take their effects from the pilot. For each language, the planned model is fitted to the pilot ratings. On every simulated study, the fixed effects and ordinal thresholds are drawn together from one draw of that fit's posterior, and the random-effect covariance is held at its posterior mean. New participants and new verb deviations are then drawn around the pilot's own verbs and their measured affectedness scores. Each reported rate is an assurance, a detection rate averaged over the effect sizes the pilot leaves plausible, which allows for the uncertainty of a single pilot's point estimate (Albers & Lakens, 2018). Every pilot language goes forward whatever its pilot showed, so the follow-up bias that Albers and Lakens describe does not arise.

Published effect sizes are subject to publication bias and tend to overstate true effects (Button et al., 2013; Gelman & Carlin, 2014), so design analyses anchored to them can be optimistic (Vasishth et al., 2018). In this report, published estimates enter only the multi-language synthesis, for the four languages the pilot does not cover, and there they carry extra uncertainty to allow for the change of protocol. Their average, -0.34 per standard deviation of affectedness, is close to the three pilots' -0.29, so they show no sign of inflation relative to the CLAPS protocol. Like the pilot data, the published data serve only to plan the study. Neither will enter the confirmatory tests, which will use only new data collected under the registered protocol.

Affectedness is a property of the verb. In the study, it will be rated in a separate norming task, in which twenty further speakers of each language rate the affectedness of the patient in each event, once for each agent-gender version. A verb's score is the mean of its two version scores, so each verb carries one affectedness score, the same across participants and sentence types. Every affectedness effect is therefore estimated from differences between verbs, and the verbs and the spread of their scores bound its precision, however many participants take part. This is the stimulus-sampling limit long recognised in psycholinguistics (Clark, 1973) and quantified for designs with crossed random effects by Westfall et al. (2014; see also Brysbaert & Stevens, 2018; Judd et al., 2017). With unlimited participants, the standard error of such an effect cannot fall below the by-verb standard deviation of that effect divided by the square root of the verbs' summed squared affectedness deviations. That floor can be read from the fitted pilot models without simulation.

The analysis model has the maximal random-effects structure the design supports (Barr et al., 2013), with correlated by-participant slopes for sentence type, affectedness and their interaction, and by-verb intercepts and sentence-type slopes. Because affectedness does not vary within a verb, a by-verb affectedness slope cannot be estimated and is omitted. Affectedness is centred and divided by twice its standard deviation within each language, and effects are reported per standard deviation of affectedness. The priors are centred at zero on the focal slopes, so they favour neither sign. The affectedness slope and the active interaction take normal priors with a standard deviation of 0.5 on the model's scale, or 0.25 per standard deviation of affectedness, and the pseudo-passive interaction

takes 0.6, or 0.30 per standard deviation. The other fixed effects take normal priors with a standard deviation of 1.5. The ordinal thresholds take Student-t priors with three degrees of freedom and a scale of 2.5, and the random-effect standard deviations Student-t priors with three degrees of freedom and a scale of 1. Their correlations take an LKJ prior with a shape of 2 (Gelman, 2006; Lewandowski et al., 2009).

In the confirmatory test, each language is first analysed on its own with this model, so it keeps its own rating thresholds, participants and verbs. A verb's effects are therefore estimated separately in each language. This suits verbs whose translations mean similar but never identical things, and across the meta-analysis languages verb effects correlate only about 0.2. The per-language posterior means and standard deviations of each interaction are then combined in a Bayesian random-effects meta-analysis, with language as the random effect. The mean across languages takes a normal prior with the standard deviation of the per-language prior, and the between-language standard deviation, τ, a half-normal prior with a scale of 0.5. One- and two-stage analyses of data from many studies agree when each study is large and both are specified alike (Burke et al., 2017), and each language here contributes 80 participants and about 72 verbs. A single model across all languages proved hard to fit (see Limitations), so it is kept as a planned robustness check, reported beside the confirmatory test but deciding neither hypothesis.

Evidence for each hypothesis is a directional Bayes factor, the posterior odds that the coefficient, or in the synthesis the mean across languages, carries the predicted sign divided by the same odds under the prior (Morey & Wagenmakers, 2014). The hypothesis is represented by the half of the zero-centred prior on the predicted side, a half-normal distribution of the kind *Nature Human Behaviour* names for representing a prediction. Its competitor, the reversed effect, is represented by the same half-normal on the other side. Under these symmetric priors, a Bayes factor of 10 corresponds to a posterior probability of .91 that the effect has the predicted sign, and a Bayes factor of 1/10 or less is equally strong evidence for the reversal. Within a language, the Bayes factor is estimated from posterior draws and from prior draws sampled alongside them. A simulated study counts as a detection when its Bayes factor reaches 10.

Each stimulus exists in two versions, one with a woman agent and one with a man agent. Each participant will see one version of each verb, so an English or Turkish session has 216 sentences, and the simulations describe that session. The versions are balanced across participants by counterbalanced lists (Appendix B).

The per-language results rest on a dedicated set of 120 simulated studies at each of 70, 80 and 100 participants per language, each scored for both hypotheses. The studies keep the pilot's verb counts, 72 in English and Turkish and 66 in Norwegian. The repository's wider grid adds 24 earlier studies at 70 and 100 participants, scored for H1 alone. Each study is fitted with four chains of 3,000 iterations at a target acceptance rate of .99 and a maximum treedepth of 12. After 1,000 warm-up iterations per chain, 8,000 post-warm-up draws remain. The confirmatory analysis will use

sixteen chains of 5,000 iterations with 2,000 of warm-up. The shorter chains add Monte-Carlo noise to each simulated Bayes factor without shifting it systematically. Each per-language rate in Table 1 carries an exact 95% interval. Over 120 studies, a rate near 90% has a Monte-Carlo standard error of about three percentage points, so differences of a few points between neighbouring sample sizes are noise.

The simulations ran on the University of Oxford's Advanced Research Computing service under R 4.4.2, with brms 2.21.0 and the cmdstanr backend on CmdStan 2.39.0, and the lockfile `renv.lock` fixes the package versions. The analysis code, the simulation settings and the aggregated results are openly available at https://github.com/pablobernabeu/CLAPS-analysis, where H1 and H2 carry the labels h1b and h2b. This report rebuilds from those results without the raw pilot data, which are not yet public.

Results

Table 1 gives, for each pilot language analysed on its own, the detection rate for each hypothesis by sample size. These per-language rates are not registered tests. They show how precisely one language's data estimate each interaction, which the confirmatory pooled test builds on.

Table 1

Detection Rates by Language and Participant Count

Language	Participants	H1	H2	Against H2
English	70	79% [71, 86]	93% [87, 97]	0% [0, 3]
	80	90% [83, 95]	95% [89, 98]	0% [0, 3]
	100	89% [82, 94]	99% [95, 100]	0% [0, 3]
Turkish	70	56% [46, 65]	1% [0, 5]	72% [63, 80]
	80	57% [47, 66]	0% [0, 3]	82% [74, 88]
	100	54% [45, 63]	0% [0, 3]	77% [68, 84]
Norwegian	70	95% [89, 98]	–	–
	80	97% [92, 99]	–	–
	100	98% [93, 99]	–	–

Note. Each row rests on 120 simulated studies. Entries are the percentage of studies reaching a directional Bayes factor of 10, with an exact 95% interval in brackets. H1 is the active-by-affectedness interaction and H2 the pseudo-passive-by-affectedness interaction, both predicted negative. The last column gives the percentage of studies in which the Bayes factor for H2 is 1/10 or less, evidence against it. The Norwegian materials include no pseudo-passive.

H1 in Each Pilot Language

At 80 participants, H1 reaches a Bayes factor of 10 in 90% [83, 95] of simulated English studies, 97% [92, 99] of Norwegian ones and 57% [47, 66] of Turkish ones. The Turkish pilot interaction is the weakest of the three, at −0.15 logits per standard deviation of affectedness against −0.44 in English and −0.27 in Norwegian.

The rates change little between 80 and 100 participants, as the verb limit predicts, so recruiting beyond 80 would add little in any language. Even with unlimited participants, and averaging over the pilot posterior, the interaction would reach a Bayes factor of 10 in about 97% of English studies, over 99% of Norwegian ones and about 65% of Turkish ones.

H2 in Each Pilot Language

In English, the pilot's pseudo-passive-by-affectedness interaction is −0.38 logits per standard deviation, with a 95% credible interval from −0.56 to −0.20 and all 8,000 posterior draws below zero. At 80 participants, H2 reaches a Bayes factor of 10 in 95% [89, 98] of simulated studies, more often than H1.

In Turkish, the pilot interaction runs the other way, at +0.16 per standard deviation, with a 95% credible interval from +0.03 to +0.28, so the affectedness slope is steeper for Turkish pseudo-passives than for passives. The simulations draw their effects from that posterior and inherit the reversal. At 80 participants, 82% of simulated Turkish studies give a Bayes factor for H2 of 1/10 or less, evidence against it, and the remaining 22 of 120 are inconclusive. By contrast, the per-language models of the meta-analysis show the negative interaction in Balinese, Indonesian and Mandarin.

The Confirmatory Pooled Test

The per-language rates above come from the three pilot languages, and 13 of the 20 planned languages have no data yet. A normal approximation to the per-language test, calibrated against the three pilots' simulated rates, gives the expected rate for such a language analysed on its own. If its effect comes from the population described below, with the spread across the measured languages, it reaches a Bayes factor of 10 in about 63% of studies on H1 and 72% on H2. Across the three spreads considered there, the rates run from 61% to 72% for H1 and from 49% to 86% for H2.

A test on the mean across languages needs more than a few languages. With the three pilots analysed together in the single model across languages, the active interaction reaches a Bayes factor of 10 in only 57% [47, 67] of 107 simulated studies at 50 to 100 participants per language. The rate shows no trend over that range. It is held down by the variation of the effect between languages, which neither verbs nor participants reduce and three languages estimate poorly. These simulations

used the structure of the published meta-analysis model, with its by-verb affectedness slope. Another 13 of the 120 planned studies, 11 of them at 100 participants, returned no result, and counting them as failures would lower the rate to 51%.

The collaborators expect at least 20 languages, 15 of them with a pseudo-passive, so the confirmatory test is simulated with 20 languages. The seven languages measured so far, the three CLAPS pilots and the four other meta-analysis languages, take part as CLAPS samples, with effects drawn from their estimates. The other 13 are drawn from a population of languages like the measured ones, with a spread estimated in turn from different sets of them. Each language's estimate is simulated at the precision that 80 participants and 72 verbs achieve, and each simulated study is analysed as registered. H2 is simulated in the same way in the 15 languages with a pseudo-passive, five of them among the seven, and Table 2 gives both results.

Table 2

Pooled Detection Rates with Twenty Languages, the Seven Measured Ones Among Them

	H1, active		H2, pseudo-passive	
New languages as different as	Median τ	Detection rate	Median τ	Detection rate
The CLAPS pilots	0.16	92%	0.42	56%
All measured languages	0.39	88%	0.45	89%
The other meta-analysis languages	0.53	74%	0.39	94%

Note. Each rate is the percentage of 4,000 simulated studies in which the directional Bayes factor of the two-stage synthesis reaches 10. Each study has 20 languages, 15 of them with a pseudo-passive, and 80 participants and 72 verbs per language. τ is the between-language standard deviation of the interaction per standard deviation of affectedness, given as its posterior median. For H2, each set keeps only its languages with a pseudo-passive: two pilots, five measured languages and three meta-analysis languages. Each meta-analysis language carries an extra standard deviation of 0.15 for the change of protocol, or 0.40 in the first H1 row, whose small spread implies that the earlier estimates carry over poorly.

The spread of the effect between languages decides the result, and the languages measured so far estimate it only roughly. For H1, the three CLAPS pilots differ little, the four other meta-analysis languages differ a great deal, and all seven together lie between. The larger spread may reflect real differences between the languages, or the earlier studies' procedures, which differed in scale length, verb count and affectedness norming. The planning value is the spread across all seven languages, since three languages are too few to estimate a spread well. On that value, the pooled test of H1 reaches a Bayes factor of 10 in 88% [87, 89] of simulated studies. The rate rises to 92% if new

languages differ as little as the pilots, and falls to 74% if they differ as much as the meta-analysis languages.

For H2, the planning value is again the spread across the measured languages, here the five with a pseudo-passive. On it, the pooled test reaches a Bayes factor of 10 in 89% [88, 90] of simulated studies. This time the pilots are the unfavourable set. The Turkish reversal offsets the English effect, so new languages drawn around the two average close to zero, and on their spread the rate is 56%.

The Sampling Plan

Every team will collect usable data from 80 participants, replacing any participant who is excluded. The number is fixed in advance because each team must state its sample size to its ethics committee. Each language will use as close to 72 verbs as its supply of suitable verbs allows. Every language enters the pooled tests whatever its pilot showed, Turkish included, since leaving a language out on the strength of its pilot would amount to choosing languages by their results. The H2 rates above already allow for the Turkish reversal. Table 3 summarises the planned analyses and their detection rates.

Table 3

Planned Analyses and Their Detection Rates

Analysis	Coefficient and direction	Test	Detection rate
H1, primary	Active by affectedness, mean across languages, negative	Directional Bayes factor of 10 in the two-stage synthesis	88% with 20 languages
H2, secondary	Pseudo-passive by affectedness, mean across languages, negative	As for H1, in the languages with a pseudo-passive	89% with the 15 of those 20 that have one
Robustness check	Both interactions in a single model across languages	Reported with its directional Bayes factors, deciding neither hypothesis	About 89% and 90%, by an analysis calibrated to behave like it

Note. The rates are at 80 participants per language, with 20 languages, on the spread across the measured languages that carry the interaction.

Nature Human Behaviour sets its 0.95 power requirement for frequentist analysis plans. For a Bayesian plan, it asks that data collection continue until the Bayes factor is "at least 10 times in favour of the experimental hypothesis over the null hypothesis (or vice versa)". Authors with limited resources may instead state a maximum feasible sample, "sufficiently large that inconclusive results at this sample size would nevertheless be an important message for the field". CLAPS takes the second route for both hypotheses. Each team stops at 80 participants, the most the collecting teams can provide, and data collection closes on a date fixed in advance. A test whose Bayes factor has reached neither 10 nor 1/10 by then is reported as inconclusive. At that sample, each language is close to the limit its verbs set, so further participants would add little. An inconclusive result from at least 20 languages and 1,600 participants would still be an important message for the field.

Limitations

At 80 participants, no simulated fit has a divergent transition or a saturated treedepth. Confirmatory fits must also reach an R-hat below 1.01 and bulk and tail effective sample sizes of at least 400. With the shorter runs used for simulation, 27 of 120 English fits meet those thresholds, as do 33 Norwegian and 12 Turkish fits. The rates above count all fits, and restricting them to the passing fits gives 81% in English, 100% in Norwegian and 50% in Turkish, each within what a random subset of that size would show. With longer chains and four times as many of them, the confirmatory fits will have six times the post-warm-up draws.

Fitted to the combined pilot data, 75,528 ratings over 210 verbs, a single model across languages estimated the interaction at about -0.24 logits per standard deviation of affectedness, with a 95% credible interval from -0.51 to $+0.15$. That is close to the average of the three pilot estimates, -0.29, but the model is much harder to fit than the per-language ones. It saturated the maximum treedepth on three quarters of its iterations, a problem longer chains do not fix. In simulations with 20 languages of 20 verbs each, no form more complex than one with sentence-type slopes for participants and verbs and uncorrelated by-language slopes converged. The robustness check's model ladder is therefore expected to end at that form.

The 20-language figures rest on an assumed population of languages. The seven measured so far are a convenience sample from five families in two macroareas, and the simulations treat the other languages as draws from a population like them. They also give each language's estimate the standard error of a normal approximation to its analysis. The calibration behind the rates for a language without a pilot widens that standard error by a factor of 1.17, and with it every rate in Table 2 falls by at most 0.4 percentage points.

References

Albers, C., & Lakens, D. (2018). When power analyses based on pilot data are biased: Inaccurate effect size estimators and follow-up bias. *Journal of Experimental Social Psychology*, *74*, 187–195. https://doi.org/10.1016/j.jesp.2017.09.004

Ambridge, B., Arnon, I., & Bekman, D. (2023). He was run-over by a bus: Passive – but not pseudo-passive – sentences are rated as more acceptable when the subject is highly affected. New data from Hebrew, and a meta-analytic synthesis across English, Balinese, Hebrew, Indonesian and Mandarin. *Glossa Psycholinguistics*, *2*(1), 11. https://doi.org/10.5070/G6011177

Ambridge, B., Bidgood, A., Pine, J. M., Rowland, C. F., & Freudenthal, D. (2016). Is passive syntax semantically constrained? Evidence from adult grammaticality judgment and comprehension studies. *Cognitive Science*, *40*(6), 1435–1459. https://doi.org/10.1111/cogs.12277

Aryawibawa, I. N., & Ambridge, B. (2018). Is syntax semantically constrained? Evidence from a grammaticality judgment study of Indonesian. *Cognitive Science*, *42*(8), 3135–3148. https://doi.org/10.1111/cogs.12697

Barr, D. J., Levy, R., Scheepers, C., & Tily, H. J. (2013). Random effects structure for confirmatory hypothesis testing: Keep it maximal. *Journal of Memory and Language*, *68*(3), 255–278. https://doi.org/10.1016/j.jml.2012.11.001

Beavers, J. (2011). On affectedness. *Natural Language & Linguistic Theory*, *29*(2), 335–370. https://doi.org/10.1007/s11049-011-9124-6

Bidgood, A., Pine, J. M., Rowland, C. F., & Ambridge, B. (2020). Syntactic representations are both abstract and semantically constrained: Evidence from children's and adults' comprehension and production/priming of the English passive. *Cognitive Science*, *44*(9), e12892. https://doi.org/10.1111/cogs.12892

Brysbaert, M., & Stevens, M. (2018). Power analysis and effect size in mixed effects models: A tutorial. *Journal of Cognition*, *1*(1), 9. https://doi.org/10.5334/joc.10

Burke, D. L., Ensor, J., & Riley, R. D. (2017). Meta-analysis using individual participant data: One-stage and two-stage approaches, and why they may differ. *Statistics in Medicine*, *36*(5), 855–875. https://doi.org/10.1002/sim.7141

Bürkner, P.-C., & Vuorre, M. (2019). Ordinal regression models in psychology: A tutorial. *Advances in Methods and Practices in Psychological Science*, *2*(1), 77–101. https://doi.org/10.1177/2515245918823199

Button, K. S., Ioannidis, J. P. A., Mokrysz, C., Nosek, B. A., Flint, J., Robinson, E. S. J., & Munafò, M. R. (2013). Power failure: Why small sample size undermines the reliability of neuroscience. *Nature Reviews Neuroscience*, *14*(5), 365–376. https://doi.org/10.1038/nrn3475

Clark, H. H. (1973). The language-as-fixed-effect fallacy: A critique of language statistics in psychological research. *Journal of Verbal Learning and Verbal Behavior*, *12*(4), 335–359. https://doi.org/10.1016/S0022-5371(73)80014-3

Darmasetiyawan, I. M. S., & Ambridge, B. (2022). Syntactic representations contain semantic information: Evidence from Balinese passives. *Collabra: Psychology*, *8*(1), 33133. https://doi.org/10.1525/collabra.33133

Dowty, D. R. (1991). Thematic proto-roles and argument selection. *Language*, *67*(3), 547–619. https://doi.org/10.1353/lan.1991.0021

Gelman, A. (2006). Prior distributions for variance parameters in hierarchical models (comment on article by Browne and Draper). *Bayesian Analysis*, *1*(3), 515–534. https://doi.org/10.1214/06-BA117A

Gelman, A., & Carlin, J. (2014). Beyond power calculations: Assessing Type S (sign) and Type M (magnitude) errors. *Perspectives on Psychological Science*, *9*(6), 641–651. https://doi.org/10.1177/1745691614551642

Judd, C. M., Westfall, J., & Kenny, D. A. (2017). Experiments with more than one random factor: Designs, analytic models, and statistical power. *Annual Review of Psychology*, *68*, 601–625. https://doi.org/10.1146/annurev-psych-122414-033702

Lewandowski, D., Kurowicka, D., & Joe, H. (2009). Generating random correlation matrices based on vines and extended onion method. *Journal of Multivariate Analysis*, *100*(9), 1989–2001. https://doi.org/10.1016/j.jmva.2009.04.008

Liddell, T. M., & Kruschke, J. K. (2018). Analyzing ordinal data with metric models: What could possibly go wrong? *Journal of Experimental Social Psychology*, *79*, 328–348. https://doi.org/10.1016/j.jesp.2018.08.009

Liu, L., & Ambridge, B. (2021). Balancing information-structure and semantic constraints on construction choice: Building a computational model of passive and passive-like constructions in Mandarin Chinese. *Cognitive Linguistics*, *32*(3), 349–388. https://doi.org/10.1515/cog-2019-0100

Morey, R. D., & Wagenmakers, E.-J. (2014). Simple relation between Bayesian order-restricted and point-null hypothesis tests. *Statistics & Probability Letters*, *92*, 121–124. https://doi.org/10.1016/j.spl.2014.05.010

Pinker, S., Lebeaux, D. S., & Frost, L. A. (1987). Productivity and constraints in the acquisition of the passive. *Cognition*, *26*(3), 195–267. https://doi.org/10.1016/S0010-0277(87)80001-X

Schönbrodt, F. D., & Wagenmakers, E.-J. (2018). Bayes factor design analysis: Planning for compelling evidence. *Psychonomic Bulletin & Review*, *25*(1), 128–142. https://doi.org/10.3758/s13423-017-1230-y

Taylor, J. E., Rousselet, G. A., Scheepers, C., & Sereno, S. C. (2023). Rating norms should be calculated from cumulative link mixed effects models. *Behavior Research Methods*, *55*(5), 2175–2196. https://doi.org/10.3758/s13428-022-01814-7

Vasishth, S., Mertzen, D., Jäger, L. A., & Gelman, A. (2018). The statistical significance filter leads to overoptimistic expectations of replicability. *Journal of Memory and Language*, *103*, 151–175. https://doi.org/10.1016/j.jml.2018.07.004

Veríssimo, J. (2021). Analysis of rating scales: A pervasive problem in bilingualism research and a solution with Bayesian ordinal models. *Bilingualism: Language and Cognition*, *24*(5), 842–848. https://doi.org/10.1017/S1366728921000316

Westfall, J., Kenny, D. A., & Judd, C. M. (2014). Statistical power and optimal design in experiments in which samples of participants respond to samples of stimuli. *Journal of Experimental Psychology: General*, *143*(5), 2020–2045. https://doi.org/10.1037/xge0000014

Appendix A

Composition of the Pilot Sample

This appendix describes the pilot sample whose ratings give the descriptive trends in Figure 1 and to which the planned model is fitted for the simulations. Table A1 gives the counts, read directly from the harmonised pilot data.

Table A1

Composition of the Harmonised Pilot Sample, by Language

Language	Participants	Verbs	Sentence types	Items	Judgements
English	70	72	3	144	30,240
Turkish	70	72	3	144	30,240
Norwegian	57	66	2	132	15,048

Note. Items are verb-by-agent-gender combinations, so the item count is twice the verb count. Every pilot participant saw both versions of every verb, whereas in the study itself each participant will see one version of each verb (Appendix B). The Norwegian materials include no pseudo-passive, hence only two sentence types. Norwegian sessions also included a synthetic passive, which the preregistered analysis excludes, so a Norwegian participant rated one and a half times as many sentences as this table counts.

Appendix B

The Agent-Gender Lists

This appendix sets out how the two agent-gender versions of each stimulus are balanced across participants, and why. Each participant will see one version of each verb, so the balance has to come from the assignment of versions to participants.

The versions need balancing because agent gender moves the interaction in English. A linear mixed model of the 1-to-7 pilot ratings, used here only to gauge the relative size of the gender term, codes agent gender as plus or minus a half. In it, the English interaction is −0.16 rating points per standard deviation of affectedness, and its gender-by-active-by-affectedness term is +0.068 (SE = 0.023). The interaction is therefore markedly stronger for man-agent items than for woman-agent ones. A design in which every participant saw only woman-agent versions, or only man-agent ones, would move the English interaction by about 21%, against 1% in Turkish and 3% in Norwegian.

Counterbalanced lists remove that shift without lengthening the session. Each list shows woman agents on half the verbs and man agents on the other half, with the same agent across a verb's sentence types, and across the lists each verb appears equally often with each agent. Each participant therefore sees both agent genders equally often, each verb is rated in each version by half the sample, and the estimand remains the slope averaged over agent gender. The protocol uses eight such lists, although two would be enough. With two, the verbs would be split into two sets of 36, or 33 in Norwegian, matched on affectedness and verb class. One list would show the first set with woman agents and the second with man agents, and the other the reverse. As long as every list shows both agent genders on matched verb sets, unequal numbers of participants across the lists leave the interaction essentially unchanged. Even so, assigning participants to the lists in turn, and replacing exclusions within the same list, keeps each verb's two versions evenly rated. Splitting the agent-gender versions across participants in this way follows Ambridge et al. (2016, Study 3) and the Hebrew study of Ambridge et al. (2023).

Both agent-gender versions should stay in the affectedness rating task, which asks for one judgement per verb and agent gender and is short. Halving it would leave each agent-specific score resting on half as many raters and would attenuate every slope estimated from those scores.